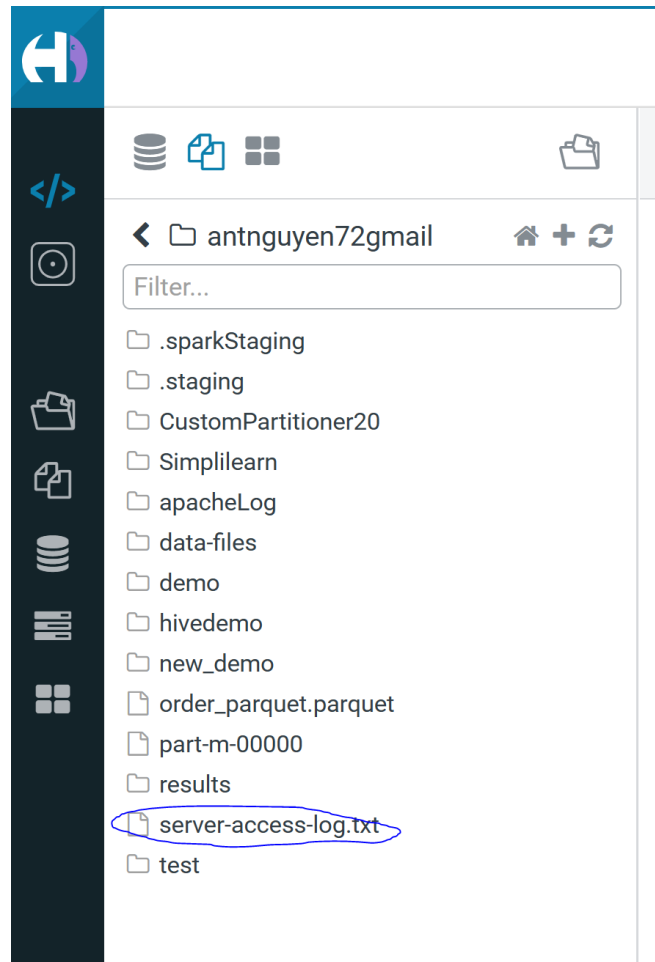


Objective:

Perform server log analysis to assist businesses in identifying and analyzing critical business errors, as well as potential customers and their domains

Steps to be performed:

Step 1: Upload the “server-access-log” file to the HDFS



Step 2: Execute the PySpark Commands following the below steps:

- Login to the web console
- To enter the PySpark console run the following command:

Command: pyspark3

I am using Jupyter Notebook

Step 3: Perform the below tasks on the uploaded dataset:

- Status code analysis:
 - Read the log file as an RDD in PySpark
- Consider the sixth element as it is “request type” and replace the “single quote” with blank
 - Convert each word into a tuple of (word,1)
 - Apply “reduceByKey” transformation to count the values

- Display the data

```
# Read the Log file as an RDD in PySpark via Spark Session and Spark Context
from pyspark.sql import SparkSession
spark = SparkSession.builder.appName('spark').getOrCreate()

log_rdd = spark.sparkContext.textFile('server-access-log.txt')

code_rdd = log_rdd.filter(lambda line : line != "")\
    .map(lambda line : (line.split(" ")[5].replace('"', ''),1))\
    .reduceByKey(lambda a,b:a+b)
# .filter for taking out empty lines
# .map for doing transformation for each line --> applying function to each line
# .reduceByKey is a narrow aggregation function that does local reduce on each
# node locally before sending the reduced values to the reduce layer, causing
# less network overhead

# print result
for line in code_rdd.collect():
    print(line)

('GET', 41075)
('HEAD', 307)
('POST', 56995)
('OPTIONS', 1)
```

- Arrange the result in descending order and display the result

```
[70]: # sort by descending order
ranked_rdd = code_rdd.sortBy(lambda x : x[1],ascending = False)
for line in ranked_rdd.collect():
    print(line)

('POST', 56995)
('GET', 41075)
('HEAD', 307)
('OPTIONS', 1)
```

- Identify the top 10 frequent visitors of the website and show the result in the RDD

```

79]: # Lets assume an IP address counts as a visitor
visitor_rdd = log_rdd.filter(lambda x : x != "")\
    .map(lambda line : (line.split(' ')[0],1))\
    .reduceByKey(lambda a,b : a+b)\
    .sortBy(lambda x : x[1], ascending = False)

# filter function eliminates any empty lines
# map function transform the row into a tuple (keyIWantToKeep,1) which are later used by
# reduceByKey function to count up all instances of the IP address
# sortBy is finally used to sort via descending order

80]: for line in zip(visitor_rdd.collect(),range(10)):
    print(line)

(('193.106.31.130', 43423), 0)
(('173.255.176.5', 5220), 1)
(('178.44.47.170', 2824), 2)
(('51.210.183.78', 2684), 3)
(('45.15.143.155', 1927), 4)
(('45.144.0.179', 946), 5)
(('176.222.58.254', 934), 6)
(('45.132.207.154', 890), 7)
(('45.153.227.55', 888), 8)
(('45.138.4.22', 880), 9)

```

- Identify the top 10 missing (does not exist) URLs using these steps:
 - Read the log file as an RDD in PySpark
- Identify the URLs for which the server is returning the 404-request code and display the data

```

: missing_rdd = log_rdd.map(lambda line : line.split(' ')[8:11])\
    .filter(lambda line : line[0] == '404')\
    .map(lambda line : (line[2].replace(' ',''),1))\
    .reduceByKey(lambda a,b : a+b)\
    .sortBy(lambda line : line[1],ascending=False)

"""
    After some data exploratory analysis, I was able to determine that my two needed values are
    at the 8th and 11th position. Because of this, I used map() to slice the important piece
    Then, I use filter to collect entry points with 404 request code
    Then, I use map() again to remove extras on the url string
    Finally, I used reduceByKey() to count the number of instances of each url, then sort desc
"""

: # I could've filtered out '-', but I decided to keep it in here because
# it reflects the data most accurately ... the most missing url is ... a missing url

for line in zip(range(11),missing_rdd.collect()):
    print(line)

(0, ('-', 3070))
(1, ('http://www.almhuetten-raith.at', 609))
(2, ('http://www.almhuetten-raith.at/', 447))
(3, ('http://www.almhuetten-raith.at/apache-log/access.log', 398))
(4, ('http://www.almhuetten-raith.at/apache-log/', 183))
(5, ('http://almhuetten-raith.at/', 153))
(6, ('http://www.almhuetten-raith.at/index.php?option=com_phocagallery&view=category&id=1&Itemid=53', 90))
(7, ('http://www.almhuetten-raith.at/index.php?option=com_content&view=article&id=49&Itemid=55', 68))
(8, ('http://www.almhuetten-raith.at/index.php?option=com_content&view=article&id=50&Itemid=56', 53))
(9, ('http://www.almhuetten-raith.at/robots.txt', 51))
(10, ('http://www.almhuetten-raith.at/index.php?option=com_content&view=article&id=46&Itemid=54', 29))

```

- Identify the traffic (total number of HTTP requests received per day) using the below steps:
 - Read the log file as an RDD in PySpark
 - Fetch the DateTime string and replace "[" with blank
 - Get the date string from the DateTime
 - Identify HTTP requests using the map function
 - Display the data

```

]: traffic_rdd = log_rdd.filter(lambda line : line.split(' ')[7].split('/')[0] == 'HTTP')\
    .map(lambda line : (line.split(' ')[3].replace(',', '').split(':')[0],1))\
    .reduceByKey(lambda a,b:a+b)\
    .sortBy(lambda line:line[1],ascending=False)

traffic_rdd.collect()

]: [('28/Dec/2020', 7478),
 ('25/Dec/2020', 5644),
 ('18/Jan/2021', 4988),
 ('11/Jan/2021', 4283),
 ('08/Jan/2021', 4056),
 ('21/Dec/2020', 3982),
 ('23/Dec/2020', 3856),
 ('20/Dec/2020', 3698),
 ('22/Dec/2020', 3645),
 ('24/Dec/2020', 3607),
 ('07/Jan/2021', 3098),
 ('29/Dec/2020', 2919),
 ('09/Jan/2021', 2805),
 ('04/Jan/2021', 2788),
 ('17/Jan/2021', 2498),
 ('13/Jan/2021', 2475),
 ('30/Dec/2020', 2389),
 ('06/Jan/2021', 2386),
 ('03/Jan/2021', 2379),
 ('16/Jan/2021', 2328),
 ('10/Jan/2021', 2313),
 ('19/Jan/2021', 2302),
 ('12/Jan/2021', 2300),
 ('26/Dec/2020', 2269),
 ('15/Jan/2021', 2227),
 ('20/Jan/2021', 2204),
 ('27/Dec/2020', 2181),
 ('01/Jan/2021', 2165),
 ('31/Dec/2020', 2067),
 ('05/Jan/2021', 2017),
 ('14/Jan/2021', 1954),
 ('02/Jan/2021', 1942),
 ('19/Dec/2020', 1135)]

```

- Identify the top 10 endpoints that transfer maximum content in megabytes and display the data

```

max_rdd = log_rdd.filter(lambda line : line.split(' ')[9] != '-')\
    .map(lambda line : int(line.split(' ')[9]))\
    .sortBy(lambda line:line,ascending=False)

for line in zip(max_rdd.map(lambda x:str(x/1000000)+" Mb").collect(),range(10)):
    print(line)

('19.734268 Mb', 0)
('19.733582 Mb', 1)
('19.733209 Mb', 2)
('19.732606 Mb', 3)
('19.689319 Mb', 4)
('19.675282 Mb', 5)
('19.674632 Mb', 6)
('19.666675 Mb', 7)
('19.666118 Mb', 8)
('19.655908 Mb', 9)

```